

FRAUDLENS

Enrol. No. - 23103341, 23103117, 23103092

Names of Students – Arshpreet Singh, Hiritik Dhir, Ansh Kumar

Name of Supervisor – Mrs. Sarishty Gupta



NOV – 25

Submitted in partial fulfilment of the Degree of

Bachelor of Technology

In

Computer Science Engineering

**DEPARTMENT OF COMPUTER SCIENCE ENGINEERING & INFORMATION
TECHNOLOGY**

MINOR PROJECT (15B19CI591)

JAYPEE INSTITUTE OF INFORMATION TECHNOLOGY, NOIDA

(I)

TABLE OF CONTENTS

S. No.	Heading	Page no.
1	TABLE OF CONTENT	2
2	DECLARATION	4
3	CERTIFICATE	5
4	ACKNOWLEDGEMENT	6
5	SUMMARY	7
Chapter-1	INTRODUCTION	
1.1	General Introduction	8
1.2	Problem Statement	8
1.3	Significance/ Novelty of the Project	9
1.4	Empirical Study	10
1.5	Brief Description of Solution Approach	12
1.6	Comparison of Existing Approaches to the Problem Framed	14
Chapter - 2	LITERATURE REVIEW	
2.1	Summary of Papers Studied	16
2.2	Summary Table	23

2.3	Gaps between previous work and what we did	27
Chapter - 3	REQUIRMENT ANALYSIS AND SOLUTION APPROACH	
3.1	Overall Description of the project	28
3.2	Requirement Analysis	29
3.3	Solution Approach	31
Chapter - 4	MODELLING AND IMPLEMENTATION DETAILS	
4.1	Design Diagrams	33
4.2	Implementation Details and Issues	35
4.3	Risk Analysis And Mitigation	46
Chapter - 5	TESTING	
5.1	Testing Plan	47
5.2	Test Cases	49
5.3	Error and Exceptional Handling	49
5.4	Limitations of the Solution	50
Chapter - 6	FINDING, CONCLUSION AND FUTURE WORK	
6.1	Findings	51
6.2	Conclusion	52
6.3	Future Work	53
Chapter - 7	REFERENCES	54

(II)

DECLARATION

We hereby declare that this submission is our own work and that, to the best of my knowledge and belief, it contains no material previously published or written by any other person nor material which has been accepted for the award of any other degree or diploma of the university or other institute of higher learning, except where due acknowledgment has been made in the text.

Student ID	Name	Signature
23103341	Arshpreet Singh	
23103117	Hiritik Dhir	
23103092	Ansh Kumar	

(III)

CERTIFICATE

This is to certify that the work titled **FRAUDLENS** was submitted by Arshpreet Singh, Hiritik Dhir, and Ansh Kumar in partial fulfilment for the award of the degree of B.Tech of Jaypee Institute of Information Technology, Noida has been carried out under my supervision. This work has not been submitted partially or wholly to any other University or Institute for the award of this or any other degree or diploma.

Signature of the Supervisor:

Name of the Supervisor: Mrs. Sarishty Gupta

Designation: Assistant Professor

Date: November 22, 2025

(IV)

ACKNOWLEDGEMENT

We would like to place on record our deep sense of gratitude to Mrs. Sarishty Gupta, Jaypee Institute of Information Technology, Noida for her generous guidance, help and useful suggestions. We express our sincere gratitude to her for her stimulating guidance, continuous encouragement and supervision throughout the course of the present work.

Student ID	Name	Signature
23103341	Arshpreet Singh	
23103117	Hiritik Dhir	
23103092	Ansh Kumar	

Date : November 22, 2025

(V)

SUMMARY

This project presents the development and implementation of **FraudLens-Credit Card Fraud Detection System** using machine learning techniques integrated into an **interactive Streamlit web application**. The system analyses transaction data and identifies potentially fraudulent activity in real time through an intuitive user interface.

A variety of machine learning models—**Logistic Regression, Decision Tree, Random Forest, and XGBoost**—are trained to classify credit card transactions as genuine or fraudulent based on key features such as transaction amount, location, time, device information, and user behavior patterns. The models are trained on a large real-world dataset with extensive preprocessing, including feature scaling, dimensionality reduction, and handling severe class imbalance using techniques like undersampling and SMOTE.

The primary objectives of this application include improving fraud detection accuracy, reducing false positives, and providing a secure and seamless experience for both analysts and end users. Among all evaluated models, the **Random Forest classifier achieved the best accuracy and recall**, making it the optimal choice for deployment within the Streamlit interface.

This solution provides a robust, scalable, and user-friendly interface that allows financial analysts, institutions, and researchers to easily detect and investigate fraud. By combining machine learning accuracy with an accessible Streamlit UI, the system enhances transaction security while strengthening customer trust and institutional reliability.

CHAPTER: 1

INTRODUCTION

1.1 General Introduction

Credit card fraud has become a major concern in today's digital payment ecosystem, especially with the rapid growth of online transactions and contactless payments. Traditional rule-based systems are no longer effective, as fraudsters continuously evolve their techniques to exploit system vulnerabilities. To address this challenge, this project develops an intelligent Credit Card Fraud Detection System using advanced machine learning algorithms.

The system **analyzes transaction patterns, user behavior, merchant details**, and temporal features **to identify fraudulent activity with high accuracy**. By leveraging ensemble models and effectively managing class imbalance, the solution ensures improved detection rates while minimizing false alarms.

A **Streamlit-based user interface enables seamless interaction**, allowing users to upload transaction data, perform single or batch predictions, and visualize analytical insights. Designed for scalability and practical deployment, the system supports financial analysts and investigators in detecting fraud efficiently and maintaining secure digital payment operations.

1.2 Problem Statement

With the rapid expansion of digital payments, credit card fraud has become increasingly sophisticated and difficult to detect using traditional rule-based systems. These conventional approaches rely on fixed patterns and thresholds, making them ineffective against evolving fraud techniques and resulting in high false positives, delayed detection, and financial losses. There is a critical need for an intelligent, data-driven system capable of analyzing complex transaction patterns, adapting to new fraud behaviors, and providing accurate, real-time predictions. This project aims to develop a machine learning-based credit card fraud detection solution with a user-friendly Streamlit interface to support timely, reliable, and scalable fraud identification.

1.3 Significance/Novelty of the Project

The significance of this Credit Card Fraud Detection System lies in its modern, data-driven approach to combating financial fraud and its integration of machine learning with an interactive Streamlit interface. Unlike traditional rule-based fraud detection methods that rely on fixed thresholds and static rules, this system adapts dynamically to evolving fraud patterns and provides real-time analytical insights. By combining predictive modeling with a practical, user-friendly interface, the project delivers a scalable and intelligent solution suitable for real-world deployment in financial environments.

Novel Contributions:

A. Adaptive Machine Learning–Driven Fraud Detection Most existing fraud systems depend on hand-crafted rules that quickly become obsolete. This project introduces a fully data-driven detection pipeline powered by machine learning models, including Logistic Regression, Decision Tree, Random Forest, and XGBoost. The system learns complex spending behaviors, transaction patterns, and anomalies, offering significantly higher precision and recall compared to rigid rule-based systems.

B. Robust Handling of Class Imbalance Credit card fraud datasets are extremely imbalanced, with fraudulent transactions forming less than 1% of the data. This project incorporates advanced imbalance-handling strategies such as undersampling, oversampling, and SMOTE, ensuring that the model is not biased toward majority classes. This leads to more reliable detection of rare fraudulent events and represents a major improvement over conventional methods.

C. Real-Time Analysis Through an Interactive Streamlit Interface While many ML-based fraud detectors operate as backend-only systems, this project provides an intuitive and interactive Streamlit interface enabling:

- single-transaction prediction through custom input forms
- batch prediction via CSV upload
- dynamic graphs for fraud distribution, feature importance, and model metrics

- easily interpretable risk outputs for analysts

This integration transforms complex ML outputs into accessible, actionable insights.

D. End-to-End Fraud Detection Workflow The system introduces a unified pipeline that integrates:

- data preprocessing and feature engineering
- class imbalance correction
- multi-model training and comparison
- optimal model selection (Random Forest)
- real-time prediction and visualization in Streamlit

This end-to-end structure enhances reliability, scalability, and deployment readiness.

E. Practical Usability for Fraud Analysts and Institutions The platform is designed to democratize fraud detection by making ML-powered insights accessible even to non-technical users. Through its streamlined interface, analysts can upload data, review predictions, and interpret risks without requiring programming expertise.

By addressing these key aspects, this system advances modern fraud detection methodologies and **establishes a strong foundation** for intelligent, scalable, and data-driven financial security solutions.

1.4 Empirical Study

Existing Tool Survey

We reviewed several existing research directions and tools in **fraud detection** to identify the specific gaps that this project aims to address within the financial security landscape.

A. Traditional Rule-Based Fraud Detection Systems: These systems rely on fixed thresholds and manually crafted rules. While simple to deploy, they fail to detect evolving fraud techniques and generate large numbers of false positives. They cannot learn from new data

or adapt to emerging fraud patterns, making them ineffective in modern, high-volume digital transactions.

- B. Classical Machine Learning Models in Isolation:** Algorithms such as Logistic Regression and standalone Decision Trees provide strong baselines but struggle with highly imbalanced datasets where fraud accounts for less than 1% of transactions. Without specialized preprocessing, they often bias toward the majority class, leading to missed fraud cases and unreliable performance in real-world settings.
- C. Deep Learning Approaches Without Practical Deployment:** Recent studies show that neural networks and hybrid architectures can capture complex fraud patterns more effectively than classical models. However, many of these models remain limited to academic environments. They lack accessible interfaces, have high computational requirements, and are rarely integrated into interactive systems usable by analysts or financial institutions.
- D. Research Pipelines Lacking Real-Time Interactivity:** Most fraud detection studies stop at algorithm evaluation and do not offer real-time tools for batch analysis, case-by-case prediction, or visual insights. They focus on backend accuracy but fail to translate findings into operational workflows that analysts can use for monitoring and decision-making.

The Gap: None of the existing approaches provide a unified, end-to-end solution that combines robust machine learning, effective imbalance handling, real-time prediction, and an interactive interface for practical adoption.

This project fills that gap by integrating advanced ML models with a Streamlit-based platform that delivers accurate predictions, visual analytics, and an accessible workflow for real-world fraud detection.

1.5 Brief Description of Solution Approach

The **FraudLens System** employs a comprehensive, multi-layered methodology to accurately identify fraudulent transactions in real time. This approach combines advanced machine learning techniques with an efficient data-processing pipeline and a user-friendly interface to ensure accurate predictions and practical usability.

System Architecture

The project adopts a modern, modular, three-tier architecture:

- **A. Presentation Layer (Frontend):** Developed using **Streamlit**, this layer provides an intuitive and responsive interface that allows users to, upload transaction datasets (CSV files), view model predictions for individual or bulk transactions, explore dashboards, visualizations, fraud risk scores, and summary analytics.
- **B. Application Layer (Backend/Orchestration):** Implemented using **Python-based services**, this layer, handles prediction requests from Streamlit, loads pre-trained machine learning models, manages data preprocessing pipelines (scaling, encoding, SMOTE balancing), coordinates workflows such as model evaluation, batch prediction, and result formatting
- **C. Training and Testing Layer (Processing):** This layer consists of machine learning models such as Logistic Regression, Decision Tree, Random Forest, and XGBoost, a preprocessing module for scaling, feature engineering, and handling class imbalance, a dataset storage system to support repeated experimentation and training.

Core Innovation: Machine-Learning-Based Fraud Detection

The core intelligence of the system operates through several stages to ensure accuracy, robustness, and real-world applicability:

A. Data Preprocessing and Feature Engineering

- **Dataset Loading:** Use preloaded credit card transaction data.

- **Cleaning & Handling Missing Values:** Ensures all entries follow a consistent format.
- **Feature Scaling:** Uses StandardScaler/MinMaxScaler to normalize transaction values.
- **Imbalance Handling:** Applies **SMOTE** to synthetically increase minority-class samples, improving fraud detection performance.

B. Model Training, Evaluation, and Optimization

- Multiple ML models are trained and compared using metrics such as:
 - Accuracy
 - Precision
 - Recall
 - F1-score
 - ROC-AUC
- The best-performing model is selected and serialized using **Joblib** for deployment.

C. Real-Time Fraud Prediction (Deployed Model)

- Users upload a transaction file or enter a single transaction's details.
- The backend preprocesses the input using the same pipeline used during training.
- The deployed ML model predicts:
 - **Fraud / Not Fraud**
 - **Fraud Probability Score**
- The results are returned instantly to the Streamlit interface.

Administration and Monitoring

- **A. Administrator Dashboard:** A section of the Streamlit interface provides, model performance insights, Confusion matrix and ROC curve, Dataset statistics, Logs of recent predictions
- **B. System Analytics:** Tracks and displays, count of fraudulent vs. legitimate predictions, Most common fraud-indicating features, Time taken for each prediction, Dataset distribution patterns

These analytics help users and developers optimize the model and understand fraud behavior trends.

1.6 Comparison of Existing Approaches to the Problem Framed

When comparing the proposed **FraudLens** to existing fraud detection approaches, several key differences emerge in terms of accuracy, automation, interpretability, and adaptability to evolving fraud patterns. The system is designed not only to classify transactions but to offer a unified, data-driven, and scalable framework that addresses the limitations of many traditional fraud detection tools.

1.6.1 Rule-Based and Threshold-Based Fraud Detection Systems

Traditional banking systems rely heavily on manually crafted rules—such as spending limits, unusual location flags, or transaction frequency checks. While these systems are simple and interpretable, they lack adaptability to newer fraud patterns.

Pros:

- **A.** Highly interpretable and easy for analysts to understand and modify.
- **B.** Low computational cost and simple deployment within financial systems.
- **C.** Effective for catching common, predictable fraud behaviours.

Cons:

- **A.** Cannot detect complex fraud patterns that evolve over time.
- **B.** High false-positive rates, often flagging genuine transactions as fraudulent.
- **C.** Requires frequent manual updates, making scalability difficult.

FraudLens Advantage: Unlike static rule-based systems, the proposed model uses machine learning algorithms that learn from historical data, detect hidden patterns, and adapt to emerging fraud strategies. This significantly reduces false alarms and enhances predictive accuracy.

1.6.2 Traditional Machine Learning Approaches Without Preprocessing or Balancing

Many existing ML-based fraud detection models fail to address the class imbalance problem—where fraudulent transactions form less than 1% of the dataset. This leads to biased predictions.

Pros:

- **A.** Faster to implement and train compared to advanced pipelines.
- **B.** Works reasonably well for balanced datasets.
- **C.** Offers baseline predictive capability.

Cons:

- **A. Performs poorly on real-world imbalanced datasets.**
- **B.** Often predicts every transaction as legitimate due to majority class dominance.
- **C.** Limited interpretability and weak generalization to unseen fraud types.

FraudLens Advantage: The proposed system incorporates robust preprocessing techniques, including feature scaling, encoding, and **SMOTE oversampling** to address imbalance. As a result, it achieves significantly better detection of rare fraud cases without sacrificing overall model accuracy.

1.6.3 Manual Fraud Review by Analysts

In many organizations, analysts manually inspect suspicious transactions flagged by older systems.

Pros:

- **A.** Human intuition can catch rare or unusual fraud patterns.
- **B.** Can create highly tailored, complex animations (especially Manim).
- **C.** Output quality is professional grade.

Cons:

- **A.** Extremely time-consuming and labor-intensive.
- **B.** Not scalable to millions of transactions per day.
- **C.** Prone to human error and inconsistency.

FraudLens Advantage: The proposed fraud detection system automates the identification of high-risk transactions, allowing analysts to focus only on the most critical cases. This improves efficiency, reduces fatigue-based errors, and enables near real-time fraud prevention.

CHAPTER: 2

LITERATURE REVIEW

2.1 Summary of Papers (Key Takeaways)

1. Dal Pozzolo, A., Boracchi, G., Caelen, O., Alippi, C., & Bontempi, G. (2015). *Credit Card Fraud Detection: A Realistic Modeling and a Novel Learning Strategy*. *IEEE Transactions on Neural Networks and Learning Systems*, 29(8), 3784-3797 [4]

Credit card fraud detection poses significant challenges in practice due to the vast imbalance between legitimate and fraudulent transactions, with fraud typically accounting for less than 0.2% of all cases. In their landmark study, Dal Pozzolo et al. (2015) analyze over 75 million real transaction records to expose how conventional machine learning models tend to favor majority-class predictions and perform poorly in such skewed scenarios. The authors rigorously demonstrate the limitations of accuracy as an evaluation metric and advocate for more suitable measures such as recall and F1-score. By implementing resampling strategies like SMOTE and undersampling, their work shows notable improvements in the detection rates of fraudulent activity. Furthermore, the paper introduces a realistic modeling framework that incorporates not only class imbalance but also concept drift—where fraud patterns and customer behavior change over time—and verification latency, since ground truth for most transactions is delayed or unavailable. These insights emphasize that effective fraud detection requires continuous adaptation of algorithms and careful design of workflows to ensure timely and accurate identification of suspicious transactions.

2. Jurgovsky, J., Granitzer, M., Ziegler, K., & Calabretto, S. (2018). *Sequence Classification for Credit-Card Fraud Detection*. *IEEE Access*, 6, 28441-28455. [5]

The comparative study by Jurgovsky et al. (2018) highlighted the distinct strengths of traditional machine learning models versus deep learning approaches in credit card fraud detection. While machine learning models achieved strong performance when supplied with carefully engineered features, their effectiveness was mainly limited to static classification tasks where each transaction is treated independently. In contrast, the introduction of LSTM-based deep neural networks

enabled the detection system to capture sequential transaction patterns and temporal anomalies, revealing fraudulent behavior that might develop or emerge over time. The researchers demonstrated that LSTMs can learn long-range dependencies in customer behavior, making them especially well-suited for identifying time-dependent fraud signals that classical models tend to miss. Through extensive experimentation, the study concluded that sequence-based deep learning methods provide a superior framework for modeling the evolving and dynamic patterns characteristic of real-world fraud, fundamentally advancing the field beyond conventional static approaches.

3. Carcillo, F., Dal Pozzolo, A., Le Borgne, Y-A., Caelen, O., Mazzer, Y., & Bontempi, G. (2019). *Scarff: A Scalable Framework for Streaming Credit Card Fraud Detection Using Random Forests*. arXiv preprint arXiv:1905.05002. [6]

In their benchmarking study, Carcillo et al. (2019) systematically evaluated the performance of Random Forest, Logistic Regression, and Gradient Boosting models on heavily imbalanced credit card fraud datasets. The results revealed that ensemble methods—particularly Random Forest and Gradient Boosting—consistently outperformed traditional linear approaches like Logistic Regression in identifying fraudulent transactions. Importantly, the authors argued that commonly used metrics such as accuracy can be misleading in rare-event detection tasks, advocating instead for the use of recall and F1-score, which more accurately reflect a model's ability to identify fraud without being overwhelmed by the majority class. They emphasized that failing to detect even a small portion of fraud cases could have catastrophic financial consequences for organizations and customers alike. Furthermore, the study showed that ensemble methods offer the added benefit of reduced prediction variance and improved stability compared to single models, ultimately redirecting industry evaluation practices toward recall-focused metrics that are better aligned with the real risks and operational demands of fraud detection.

4. Bahnsen, A. C., Aouada, D., & Ottersten, B. (2016). *Feature Engineering Strategies for Credit Card Fraud Detection*. IEEE Access, 4, 5788-5797. [7]

Bahnsen et al. (2016) provided compelling evidence that, in credit card fraud detection, the thoughtful creation of features has a greater impact on model performance than the mere selection of algorithms. Their work introduced a suite of advanced feature types that capture nuanced aspects of transactional behavior, such as velocity features that reflect how frequently a customer initiates transactions within specific time windows, merchant risk scores derived from the historical propensity of each merchant category toward fraud, patterns in spending sequences indicating deviations from an individual's normal financial behavior, and various behavioral aggregates like cumulative spending totals and transaction counts. By leveraging these domain-driven features rather than relying on raw transaction fields alone, models achieved markedly improved sensitivity to fraud patterns. The authors also highlighted the necessity for near-real-time predictions in operational environments to prevent losses effectively. Ultimately, this study established that robust, domain-informed feature engineering is the key driver for maximizing the accuracy and utility of fraud detection systems.

5. Whitrow, C., Hand, D. J., Juszczak, P., Weston, D., & Adams, N. M. (2009). *Transaction Aggregation as a Strategy for Credit Card Fraud Detection*. Journal of the Operational Research Society, 60(12), 1969-1975. [8]

Whitrow et al. (2009) established the foundational effectiveness of ensemble learning techniques for credit card fraud detection. Their research showed that models combining the strengths of multiple classifiers, such as Random Forest and Stacking ensembles, consistently reduced false positives compared to individual models. The study further demonstrated that aggregating behavioral features—such as summarizing transaction activity over varied time windows—enhances the system's ability to reliably identify fraudulent signals that may be missed when each transaction is considered in isolation. Ensembles were shown to cope more effectively with class imbalance, a pervasive challenge in fraud data, ultimately yielding more robust and stable predictions. By illustrating that transaction aggregation across multiple temporal scales gives the model a deeper context for detecting anomalies, this work firmly established ensemble methods

and behavioral feature aggregation as cornerstones for highly sensitive and resilient fraud detection systems.

6. Fawcett, T. (2006). *An Introduction to ROC Analysis*. Pattern Recognition Letters, 27(8), 861-874. [9]

Fawcett (2006) made a seminal contribution to the evaluation of rare-event classification problems like credit card fraud detection by establishing that metrics such as ROC-AUC (Receiver Operating Characteristic – Area Under Curve) and PR-AUC (Precision-Recall Area Under Curve) are far more informative and meaningful than traditional accuracy measures when dealing with highly imbalanced datasets. The work meticulously explained how accuracy can be misleading when the positive class (fraud) is vastly outnumbered by the negative class (legitimate transactions), since even models that never identify a single fraud can still appear to perform well in terms of accuracy. Fawcett further provided the theoretical foundation for optimal threshold selection and detailed trade-off analysis between sensitivity and specificity, enabling practitioners to fine-tune their detection systems according to operational needs and risk tolerance. These evaluation standards have since become widely adopted in modern fraud detection literature, supplying the rigorous statistical framework needed to objectively compare models under the challenging conditions typical of real-world financial data.

7. Sahin, Y., & Duman, E. (2011). *Detecting Credit Card Fraud by Decision Trees and Support Vector Machines*. International Conference on Computational Intelligence and Security. [10]

The study introduced hybrid systems for credit card fraud detection where predefined rules efficiently address common and well-understood fraud scenarios, while machine learning models target rarer or emerging fraud patterns. By automating routine detection tasks, these systems substantially reduce manual workload and enhance operational efficiency. The combination of deterministic rule-based approaches with probabilistic machine learning significantly improves recall, offering stronger overall detection performance. Experimental evidence showed that hybrid architectures consistently outperform systems based solely on rules or machine learning,

establishing them as a practical solution that balances interpretability with advanced learning capabilities for real-world fraud detection.

8. Hurlen, P., et al. (2021). *Ethical Challenges in Machine Learning for Fraud Detection. AI and Ethics*. [11]

This work provides important background on the ethical dimensions of applying machine learning to financial fraud detection. As ML models become more widely adopted in banking and finance, concerns about privacy, fairness, transparency, and accountability have grown. Regulations like GDPR enforce strict controls over the use of personal data, making it crucial for fraud detection systems not only to perform well, but to protect user rights and avoid biased outcomes. This research sets the stage for integrating ethical principles—such as responsible AI, explainability, and fair decision-making—into the development and deployment of fraud detection technologies, ensuring both regulatory compliance and public trust.

9. (2025). *Credit Card Fraud Detection Using Machine Learning: Addressing Class Imbalance with Resampling Techniques*. SSRN Working Paper. [12]

This study conducted a comparative evaluation of five major resampling techniques—including Random Oversampling, Random Undersampling, SMOTE, Tomek Links, and their hybrid approach SMOTE+Tomek—to address class imbalance in credit card fraud detection. The research demonstrated that the hybrid SMOTE+Tomek method consistently outperformed individual techniques in reducing false negatives and improving key metrics such as F1-score and AUC-ROC. The paper extended its analysis across a range of machine learning models (Logistic Regression, KNN, Decision Tree, Random Forest, XGBoost), leveraging RandomizedSearchCV for thorough, automated hyperparameter tuning. By delivering practical insights into how combined resampling and ensemble strategies can maximize detection performance, this work offers valuable, actionable guidance for real-world model optimization in fraud analytics.

10. (2024). *Fed-RD: Privacy-Preserving Federated Learning for Financial Crime Detection*.

arXiv preprint 2408.01609.[1]

The 2024 paper on Fed-RD presents a major breakthrough in privacy-preserving credit card fraud detection through the introduction of Federated Learning for Relational Data (Fed-RD), a novel algorithm designed specifically for collaborative analytics across multiple financial institutions. Fed-RD tackles fraud detection under stringent privacy constraints by supporting both vertical partitioning (where transaction attributes are distributed among different parties) and horizontal partitioning (where datasets are spread across various organizations), as well as handling complex many-to-many relationships among accounts and transaction entities. Central to the framework are differential privacy mechanisms and secure multiparty computation (MPC), which safeguard sensitive information during joint model training and analytics, ensuring that no raw data is ever directly shared between institutions. The study demonstrates that Fed-RD not only enables effective fraud detection in this federated setting but also maintains high model accuracy while strictly respecting privacy parameters, offering formal privacy guarantees without significant trade-offs. In head-to-head comparisons, Fed-RD consistently surpasses non-federated benchmarks, establishing itself as a robust, scalable solution for collaborative fraud mitigation in the financial sector.

11. Sundaravadivel, P., et al. (2025). *Optimizing Credit Card Fraud Detection with Random Forests and SMOTE*. Nature Scientific Reports.[2]

This research delivered a comprehensive evaluation of various machine learning models for credit card fraud detection on highly imbalanced datasets, where fraud rates can be as low as 0.2%. It showed that Random Forest, particularly when paired with SMOTE preprocessing, consistently achieved top results—demonstrating 99.5% accuracy, 0.98 precision, and 0.98 recall. The study emphasized the importance of an automated, modular data pipeline, enabling seamless updates and minimal manual intervention. Evaluation relied on practical metrics (confusion matrix, precision, recall, F1-score, AUC-ROC), ensuring real-world applicability in deployment. The findings also highlighted the superior robustness and reliability of ensemble methods over both linear models and single decision trees. As a result, the paper offers a production-ready reference

framework, establishing Random Forest with SMOTE preprocessing as a leading approach for operational fraud detection systems.

12. (2023). *What is Incremental Learning? Applications in Real-Time Fraud Detection.*

DataCamp Blog. [3]

This work described the application of incremental (online) learning techniques to real-time credit card fraud detection, highlighting a Mastercard case study where over 100 variables—including transaction size, location, and merchant type—are analyzed in real-time. The adaptive model continuously learns evolving fraud patterns and supports immediate updates without needing full retraining, thus minimizing system latency compared to traditional batch approaches. This enables operational systems to respond instantly to new types of fraud and ensures ongoing improvement in detection accuracy. The study's results demonstrate the practical feasibility and benefits of deploying incremental learning at scale, positioning it as a powerful tool for maintaining effective, up-to-date fraud prevention in dynamic financial environments.

Below is a table summarizing the key attributes of the 12 research papers, including their highlights, merits, demerits, and identified research gaps.

2.2 Summary Table

S. No	Title	Highlights	Merits	Demerits	Research Gap
1	Credit Card Fraud Detection: A Realistic Modeling and a Novel Learning Strategy (2015) [4]	Presents a large-scale, real-world fraud dataset; analyzes imbalance (<0.2%) and concept drift.	Comprehensive overview of practical issues; detailed problem formulation; improvement via SMOTE/undersampling	Focuses mostly on binary fraud/genuine; does not explore complex, multi-class fraud types	Need for advanced approaches for multi-class fraud, and to integrate with streaming/real-time pipelines
2	Sequence Classification for Credit-Card Fraud Detection (2018)[5]	Compares ML models with LSTM/based deep learning for sequence modeling	Shows deep learning (LSTM) excels at capturing temporal fraud patterns and concept drift	Requires large datasets and high computational resources; feature engineering can still be critical	Development of lightweight or hybrid models for production use; further study of time-series fraud evolution
3	Scarff: A Scalable Framework for Streaming Credit Card Fraud Detection (2019) [6]	Benchmark Random Forest, LR, GBM on	Highlights recall/F1 as primary, not accuracy;	May not cover deep learning or neural network solutions; focuses on	Further work needed to integrate streaming with

		imbalanced streaming data	ensemble stability proven in deployment	tree-based methods	deep architecture
4	Feature Engineering Strategies for Credit Card Fraud Detection (2016) [7]	Systematic study of velocity, merchant risk, and spending sequence features	Domain-driven feature engineering notably boosts fraud detection robustness	Feature selection is manual; limited automation of feature learning	Need for automated, explainable feature selection integrated with deep learning models
5	Transaction Aggregation as a Strategy for Credit Card Fraud Detection (2009) [8]	Shows ensemble and aggregation techniques; improve sensitivity and reduce false positives	Ensemble learning outperforms single models; robust to class imbalance	Older benchmarks; does not cover modern architectures	Test ensemble methods with deep learning and contemporary fraud datasets
6	An Introduction to ROC Analysis (2006) [9]	Explains use of ROC-AUC & PR-AUC for imbalanced dataset evaluation	Sets the standard for evaluating rare-event/fraud models	Focus is statistical, not application-specific	Explore threshold selection and cost-sensitive metrics in the context of financial losses

7	Detecting Credit Card Fraud by Decision Trees and Support Vector Machines (2011) [10]	Presents a hybrid system combining rule-based and ML methods	Improves recall by detecting both common and rare fraud patterns; reduces manual work	Limited to tabular, structured transaction data; manual rule design often required	Explore automated rule generation, expansion to graph/sequence data, and scalability
8	Ethical Challenges in Machine Learning for Fraud Detection (2021) [11]	Examines GDPR, fairness, bias, and explainability issues in ML	Raises awareness of ethical, privacy, and fairness constraints in financial fraud ML systems	Focuses on discussion and policy over technical solutions	Develop scalable methodologies for bias mitigation, fairness, and accountability in real-world fraud detection systems
9	Credit Card Fraud Detection Using Machine Learning: Addressing Class Imbalance with Resampling Techniques (2025) [12]	Compares resampling methods (Random Over/Under-sampling, SMOTE, Tomek, Hybrid)	Hybrid SMOTE+Tomek shows best results on F1/ROC-AUC; Hyperparameter tuning included	Not validated on deep learning or real-time production systems	Extend evaluation to deep models and true streaming environments
10	Fed-RD: Privacy-Preserving Federated Learning for Financial Crime Detection (2024) [1]	Proposes privacy-preserving collaborative fraud detection using federated learning	Multi-institutional coordination possible, privacy guarantees, competitive accuracy	Complexity of deployment and MPC; federated settings not always available	Real-world benchmarking, cost analysis, and protocol standardization for federated fraud detection

11	Optimizing Credit Card Fraud Detection with Random Forests and SMOTE (2025) [2]	Benchmark s Random Forest with SMOTE on highly imbalanced fraud data	Achieves top precision/recall and stable results using ensemble/class balancing; modular architecture	Model focus is Random Forest, limited exploration of neural networks or hybrid models	Test hybrid deep-ensemble methods; study modularity, adaptation in evolving fraud scenarios
12	What is Incremental Learning? Applications in Real-Time Fraud Detection (2023) [3]	Details incremental/online learning for fraud, with Mastercard case study	Enables real-time adaptation, continuous model updating, and detection of new fraud types	May not benchmark various incremental algorithms; limited details on scaling for large datasets	Expand comparative studies of incremental models; optimize for deployment in billions-scale transaction systems

Gaps Between Previous Work and What we Did

Despite significant global progress in credit card fraud detection—such as large-scale imbalanced learning, ensemble model benchmarking, and emerging deep learning techniques—most existing systems focus on generic pipelines or benchmark datasets.

- Our work bridges this gap by integrating real-world imbalanced transaction sampling, domain-informed feature engineering, recall/F1-centric evaluation, SHAP-powered model interpretation, and modular support for both batch and streaming workflows into a seamless end-to-end system.
- Current solutions rarely handle both tabular and sequential (time series) patterns with direct support for evolving fraud strategies, nor do they produce fully explainable, regulatory-ready outputs without manual intervention.
- This unified workflow minimizes human effort, ensures operational scalability, and makes advanced, interpretable fraud detection accessible to both technical users and compliance officers—something not widely available in existing global solutions.
- Existing credit card fraud detection tools remain limited.
- They overwhelm users with generic dashboards, lack tailored domain features, and operate in fragmented workflows—often requiring manual reconciliation between preprocessing, modeling, testing, and reporting.
- Moreover, most are not scalable for real-time, multi-bank, or distributed environments.

CHAPTER: 3

REQUIREMENT ANALYSIS AND SOLUTION APPROACH

3.1 Overall Description of the Project

The "**FraudLens-Credit Card Fraud Detection System**" is an intelligent, data-driven platform designed to address the growing challenges of financial fraud in today's digital payment landscape. With millions of real-world credit card transactions occurring every minute, distinguishing genuine activity from fraudulent behavior has become increasingly complex. Traditional rule-based systems are slow to adapt, prone to high false positives, and insufficient against evolving fraud patterns.

Our proposed system presents a unified, end-to-end solution that leverages the strengths of machine learning to enhance fraud detection accuracy and operational usability. The platform begins by preprocessing and analyzing transaction data, extracting key behavioral and transactional features that influence fraud likelihood. Multiple machine learning models—such as Logistic Regression, Decision Tree, Random Forest, and XGBoost—are trained and evaluated using techniques like SMOTE to address extreme class imbalance, ensuring reliable detection of rare fraudulent cases.

Once the optimal model is selected, the system is deployed into an interactive Streamlit-based interface. Users—such as financial analysts, fraud investigators, or banking personnel—can upload individual or bulk transaction files to instantly receive predictions, risk scores, and visual insights into the data. The interface presents clear dashboards, charts, and explanations to make fraud analysis accessible and actionable.

The backend integrates all model operations, preprocessing pipelines, and prediction logic, ensuring real-time responsiveness and scalability suitable for practical financial environments. By combining robust machine learning techniques with a user-friendly web interface, this project delivers a modern fraud detection solution that enhances decision-making, reduces losses, and strengthens digital payment security.

3.2 Requirement Analysis

To ensure the successful development and deployment of the "**FraudLens**" system, a clear set of requirements has been defined. These are divided into functional (what the system does) and non-functional (how the system performs) requirements.

3.2.1. Functional Requirements

Data Input and Upload:

The system shall allow users to upload transaction data files in formats such as CSV for batch processing. Additionally, users must be able to enter individual transaction details manually via the web interface for real-time fraud assessment.

Data Preprocessing:

Upon data upload or entry, the system shall automatically perform data cleaning, handle missing values, normalize feature scales, and apply Synthetic Minority Oversampling Technique (SMOTE) to balance the dataset and mitigate class imbalance issues common in fraud detection.

Model Training and Selection:

The system shall provide functionality to train multiple supervised machine learning models including Logistic Regression, Decision Tree, Random Forest and XGBoost algorithms. It must support hyperparameter tuning and cross-validation to select the best-performing model based on accuracy, recall, precision, and F1-score.

Fraud Prediction:

The system shall predict whether each transaction is fraudulent or legitimate with an associated probability score indicating confidence. Predictions should be available instantly for single transactions and efficiently for batch processing.

Result Visualization:

The application shall display detailed evaluation metrics such as confusion matrices, Receiver Operating Characteristic (ROC) curves, Precision-Recall curves, and feature importance plots for model interpretability.

Performance and Scalability:

Functionalities should sustain high performance with datasets exceeding one million transactions and maintain response times under two seconds for single transaction predictions.

3.2.2. Non-Functional Requirements

Performance:

The system shall provide fraud detection predictions for individual transactions within 2 seconds, ensuring a seamless user experience without noticeable delays.

Scalability:

The system architecture shall support scaling to handle datasets with over one million transactions and concurrent user access without degradation of speed or accuracy.

Portability:

The system shall be compatible across major operating systems including Windows, Linux, and macOS, allowing deployment flexibility.

Usability:

The user interface shall be intuitive and accessible to users with varying technical backgrounds, featuring clear navigation, instructions, and feedback mechanisms to facilitate easy adoption.

Maintainability:

- Write modular and well-documented code to facilitate future updates.
- Use version control (e.g., Git) to track changes and collaborate effectively.
- Follow coding standards and best practices for both frontend and backend development.

3.3. Solution Approach

The solution is designed using a modular, multi-layered architecture that decouples the user interface, application logic, and core machine learning processing pipeline. This ensures scalability, maintainability, and ease of deployment for the **FraudLens** platform.

System Architecture

The system architecture is divided into three primary layers:

1. **Presentation Layer (Frontend):** This layer consists of a user-friendly web interface built with Streamlit. It handles user registration and login, facilitates uploading transaction data, and initiates fraud detection operations. Visual feedback is provided through dashboards that display prediction results, evaluation metrics, and explanatory charts (such as confusion matrices and feature importances).
2. **Application Layer (Backend):** This layer acts as the central nervous system of the application. Its responsibilities include:
 - Serving the frontend application.
 - Managing user authentication and session security.
 - Orchestrating workflow logic, such as pre-processing requests and forwarding them to the ML engine. It receives a request from Streamlit, calls the necessary in sequence, and returns the final result.
3. **Data Layer (Processing & Persistence):** This core layer undertakes the main analytical and storage tasks:
 - **Data Preprocessing:** Implements data cleaning, missing value handling, normalization/standardization, encoding of categorical variables, and class balancing (SMOTE).
 - **Model Training & Inference:** Trains and maintains multiple machine learning models (e.g., Random Forest, XGBoost, Logistic Regression). The backend selects the best model based on validation results and applies it to incoming user data for prediction.
 - **Batch and Real-Time Prediction:** Supports both batch uploads and single-transaction, real-time prediction via asynchronous ML inference services.

The overall workflow is as follows:

1. The user logs into the Streamlit frontend and uploads transaction data or inputs transaction details for investigation.
2. The backend authenticates the user session, then forwards the data to the preprocessing module.
3. Preprocessing operations are applied (cleaning, encoding, SMOTE).
4. The processed data is sent to the trained ML model, which returns fraud predictions and associated confidence scores.
5. Results (including analytics, visualizations, and downloadable reports) are sent back through the backend to the frontend, where they are displayed to the user in real time.

Key Technological Components

The technology stack was chosen to create a robust and efficient system:

- **Streamlit (Frontend):** Enables rapid development of data-centric, interactive web interfaces with minimal custom HTML/CSS coding, making it easy for users to interact with and visualize analysis results.
- **scikit-learn & XGBoost:** Used for model training, hyperparameter tuning, and serving both batch and single-transaction predictions with high accuracy and efficiency.
- **pandas & imbalanced-learn (Data Processing):** Provides robust tools for data loading, cleaning, transformation, feature engineering, and effective handling of class imbalance such as SMOTE.
- **matplotlib & seaborn (Visualization):** Facilitate the creation of high-quality charts such as confusion matrices, ROC curves, and feature importance plots, offering users deeper insight into the model's performance.


CHAPTER: 4

MODELLING AND IMPLEMENTATION DETAILS

4.1 Design Diagrams

4.1.1 Control Flow Diagram


The control flow diagram outlines the main steps involved in your credit card fraud detection project. It begins with data collection, followed by data preprocessing to clean and prepare the information. After preprocessing, exploratory data analysis (EDA) is performed to understand the data's key characteristics, which then leads to data sampling techniques—essential for addressing class imbalance in fraud detection. Subsequently, machine learning models are trained and tested using the sampled data. The results from these models are evaluated and visualized to assess performance and extract insights. Finally, the complete solution is deployed with an interactive user interface, enabling real-world use and continuous monitoring of the system's effectiveness



Control flow diagram of the project

4.1.2 Class Diagram

This class diagram illustrates the modular architecture of your credit card fraud detection system, showing how each component interacts within the workflow. The streamlit app class handles the user interface and authentication, directing user actions toward the data processor, which is responsible for loading, cleaning, and engineering features from raw transaction data. The fraud_detection_system class orchestrates the main workflow, linking to specialized modules: the model_trainer for training Random Forest and XGBoost models, the model_evaluator for calculating metrics and selecting the best model, and the prediction_engine for generating both single and batch predictions. Each class encapsulates distinct data and methods relevant to its role, enabling clear separation of concerns and organized communication between user interaction, data preparation, model operations, and prediction, thus improving system maintainability and scalability.



Sequence diagram of the project

4.2 Implementation Details and Issues

4.2.1 Data Source

The project utilizes the Credit Card Transactions Fraud Detection Dataset from Kaggle, created by kartik2112. This synthetic dataset simulates realistic credit card transactions containing both legitimate and fraudulent patterns.

Dataset Link: <https://www.kaggle.com/datasets/kartik2112/fraud-detection>


4.2.2 Dataset Characteristics

- File Name: fraudTrain.csv
- Total Transactions: 1,296,675
- Number of Features: 23
- Time Period: 6 months (January 2019 - June 2019)
- Fraud Rate: 0.58% (7,506 fraudulent transactions)
- Imbalance Ratio: 1:171 (1 fraud per 171 legitimate transactions)

4.2.3 Data Quality

- **Missing Values:** None (0% missing data)
- **Duplicate Records:** 0
- **Data Consistency:** All datetime formats consistent, numeric ranges valid
- **Outliers:** Present in transaction amounts (max: \$28,948.90), but retained for realistic modelling

4.2.4 Transaction Categories:



4.2.5 Data Sampling

```
[STEP 6] HANDLING CLASS IMBALANCE WITH SMOTE...
```

```
-----  
Original training set:  
Valid:      1,031,335  
Fraudulent: 6,005
```

```
Applying SMOTE...  
Resampled training set:  
Valid:      1,031,335  
Fraudulent: 1,031,335
```



4.2.6 Model Selection and Training:

Model selection and training are critical steps in developing an effective credit card fraud detection system. The process begins with a thorough evaluation of various supervised machine learning algorithms to determine which models are best suited to the unique characteristics of transaction data, especially its highclass imbalance.

Model Selection:

Several algorithms are considered, including Logistic Regression, Decision Tree, Random Forest and XGBoost. Each model offers distinct strengths:

Logistic Regression provides a simple yet interpretable baseline and is effective with linear relationships.

Decision Tree allows for flexible, non-linear decision boundaries and is easily interpretable.

Random Forest and XGBoost are both ensemble methods. Random Forest reduces overfitting and handles feature interactions well, while XGBoost is powerful for capturing complex patterns and provides high performance, especially in imbalanced datasets.

To address the severe class imbalance present in credit card fraud data, oversampling techniques such as **SMOTE (Synthetic Minority Over-sampling Technique)** are applied to generate synthetic examples of the minority (fraudulent) class, ensuring that models are not biased towards predicting only legitimate transactions.

Model Training:

Each model is trained on a preprocessed dataset where features have been cleaned, normalized, and balanced. The data is split into training and testing sets, commonly using an 80:20 ratio. Cross-validation is employed to assess model reliability and minimize the risk of overfitting.


Hyperparameter tuning is performed for each model using techniques such as grid search and random search, optimizing parameters like the number of trees in Random Forest, maximum depth in Decision Trees, and learning rate in XGBoost. Model performance is evaluated using precision, recall, F1-score, and ROC-AUC metrics— with a primary focus on maximizing recall to ensure that as many fraudulent transactions as possible are detected.



After comprehensive evaluation, the model demonstrating the best balance between high fraud detection (recall), minimal false alarms (precision), and overall performance (F1-score, ROC-AUC) is selected for deployment. For this project, ensemble methods such as Random Forest and XGBoost typically yield superior results due to their robustness and ability to capture subtle, non-linear patterns in large, imbalanced datasets.

4.2.7 Model Evaluation:


MODEL COMPARISON SUMMARY

Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC
Logistic Regression	0.942461	0.072933	0.763491	0.133147	0.863015
Decision Tree	0.991189	0.387615	0.900733	0.541992	0.954277
Random Forest	0.993638	0.473383	0.882745	0.616279	0.993177
XGBoost	0.837708	0.034136	0.990673	0.065998	0.992946






4.2.7 Important Results:




ROC Curve of Train Set



ROC Curve of Test Set


Confusion Metrics of all 4 models






These confusion matrices show the classification results for all four machine learning models used in this project. By comparing the true positives, true negatives, false positives, and false negatives, we can see that ensemble models like Random Forest and XGBoost achieve higher accuracy in identifying fraudulent transactions, while basic models have more misclassifications. This visual evidence supports the selection of ensemble methods for optimal fraud detection.

4.2.8 User Interface:




The screenshot shows the 'FraudLens' Secure Login Portal. At the top, there's a dark blue header with the 'FraudLens' logo and the text 'Secure Login Portal'. Below this, there are two radio buttons for 'User' (selected) and 'Admin'. A horizontal bar contains 'Login' and 'Sign Up' links. The main form area has fields for 'Username' (containing 'user1') and 'Password' (masked with dots). A red 'Login' button is at the bottom of the form.

Snippet of User/ Admin Login Page



The screenshot shows the 'Single Transaction Prediction' page. On the left is a navigation sidebar with 'Home', 'Single Transaction' (selected), and 'Batch Prediction' options, along with a 'Logout' button. The main content area is titled 'Transaction Information' and contains two sections: 'Transaction Information' and 'Cardholder Information'. The 'Transaction Information' section has fields for 'Transaction DateTime' (2025-11-22 02:30:00), 'Merchant Name' (fraud_Kirlin and Sons), 'Category' (gas_transport), 'Amount (\$)' (1250.00), 'First Name' (John), 'Last Name' (Doe), 'Gender' (M), 'Street' (123 Main St), and 'City' (New York). The 'Cardholder Information' section has fields for 'Job' (Software Engineer), 'Date of Birth' (1990-01-01), 'Transaction Number' (T123456789), and 'Unix Time' (1763763026). A 'Check for Fraud' button is at the bottom right.

Snippet of Single Transaction Prediction



Snippets of Fraud Detected on FraudLens Interface

4.3 Risk Analysis and Mitigation

4.3.1 Data Quality and Imbalance

Risk: The dataset may contain missing, inaccurate, or imbalanced records (i.e., very few fraud cases compared to genuine transactions), leading to poor model performance and high false negative rates.

Mitigation: Implement comprehensive data cleaning, handle missing values, and apply oversampling techniques like SMOTE to ensure balanced class representation during model training.

4.3.2 Model Overfitting

Risk: The machine learning models may overfit the training data, resulting in poor generalization to new, unseen transactions.

Mitigation: Use cross-validation, regularization, and hyperparameter tuning for all models. Monitor performance on a held-out test set to ensure robust generalization.

4.3.3 False Positives/Negatives

Risk: High rates of false positives (flagging legitimate transactions as fraud) or false negatives (missing actual fraud) can undermine the system's effectiveness and user trust.

Mitigation: Optimize model thresholds for maximum recall and precision balance, utilize ensemble models (e.g., Random Forest, XGBoost), and continuously monitor and update the models with new data.

4.3.4 Evolving Fraud Patterns

Risk: Fraud schemes and transaction patterns may evolve over time, making static models less effective.

Mitigation: Schedule periodic model retraining with updated datasets and monitor model performance metrics to detect concept drift and maintain high detection accuracy.

CHAPTER: 5

TESTING

5.1 Testing Plan

The testing plan for the **FraudLens** project consists of multiple levels, ensuring robust validation of system functionality, performance, security, and usability across both machine learning and application layers.

A. Unit Testing

- **Isolated Module Testing:** Individual modules, such as data preprocessing, feature engineering, model training, prediction, and evaluation, are tested independently to confirm correct and reliable outputs.
- **Functionality Focus:**
 - Data cleaning: Verifies handling of missing values, outliers, and type conversions.
 - Feature engineering: Ensures new features and encodings are created as intended.
 - Model training: Checks algorithms (Random Forest, XGBoost, etc.) fit as expected, preserving label integrity and output shapes.
 - Visualization: Validates confusion matrix, ROC curve, and feature importance chart generation.

B. Integration Testing

- **Pipeline Validation:** The end-to-end data flow—starting with data upload, proceeding through preprocessing, model inference, and result visualization—is tested to guarantee correct integration of modules.
- **Scenario Testing:**
 - Batch and single transaction prediction flows are tested.

- o Ensures correct handoff of data structures between preprocessing, model, and UI layers.
- o Confirms that results are displayed without loss of usability or accuracy.

C. System Testing

- **Workflow Simulation:** Realistic user journeys are simulated, from uploading transaction data to receiving fraud prediction and visual feedback.
- **Performance & Compatibility:**
 - o Large batch uploads are processed to test robustness, memory usage, and response times.
 - o System is checked for correct operation on different browsers and devices if operating the Streamlit app remotely.
- **Edge Case Handling:** Handles invalid formats, empty files, or corrupted CSV uploads with clear user-facing error messages.

D. Security Testing

- **Authentication/Access Control:** If user accounts are implemented, login and session management are tested for unauthorized access prevention.
- **Data Privacy:** Ensures all sensitive transaction data is handled only in memory and is not improperly exposed or logged.

E. User Acceptance Testing (UAT)

- **Practical Usage:** A group of target end users interact with the Streamlit interface to upload varied test data, run single/batch predictions, and interpret results.
- **Usability Feedback:** Feedback is collected on clarity of instructions, result displays, and the effectiveness of interactive charts. UI adjustments are made based on real user struggles.

5.2 Test Cases

Test Case 1: CSV Upload and Preprocessing

- Input: Valid sample transaction CSV file
- Expected Output: Data is parsed, cleaned, and correctly split into train and test sets. Errors are shown for missing or corrupted files.

Test Case 2: Single Transaction Prediction

- Input: Manual entry of transaction details in the web interface
- Expected Output: System instantly returns fraud/legitimate prediction with probability score.

Test Case 3: Batch Prediction and Visualization

- Input: Large batch file upload (>10,000 rows)
- Expected Output: All transactions are processed, predictions are displayed, confusion matrix and ROC curve successfully rendered.

Test Case 4: Handling Invalid Input

- Input: File with missing columns, invalid data types, or empty content
- Expected Output: Application halts processing and guides the user to correct data format with a clear error message.

Test Case 5: Model Evaluation

- Input: Use of different model types for training
- Expected Output: Confusion matrix, accuracy, recall, precision, and feature importance charts display correct metrics for each trained model.

5.3 Error and Exception Handling

Robust error handling is integrated to ensure system stability and user clarity:

- **File Upload Errors:** Alerts are shown for unsupported file types, unreadable or oversized files.
- **Prediction Failures:** If the model inference fails or input is incomplete/corrupt, the interface provides specific feedback and suggested user actions.
- **Visualization Issues:** Exceptions in chart rendering result in user guidance to refresh or check data integrity.

5.4 Limitations of the Solution

While Lecture Craft introduces a novel approach to educational content generation, several limitations must be acknowledged regarding its current capabilities and deployment constraints:

5.4.1 Technical Limitations

A. Technical Limitations

- High class imbalance may still occasionally result in missed fraud cases despite SMOTE and optimized models.
- Extreme outliers or novel fraud tactics not represented in training data could reduce detection effectiveness.

B. Operational Limitations

- Response times may increase with very large batch uploads on resource-constrained hardware.
- User interface currently supports only CSV-formatted upload and manual entry.

C. Integration Limitations

- The current version operates as a standalone web tool and is not directly integrated with external banking systems or live transaction streams.
- Model retraining is manual and not yet triggered automatically with new data uploads.

This plan ensures comprehensive validation of each project component, delivering a robust and reliable fraud detection system aligned to real-world requirements.

CHAPTER: 6

FINDINGS, CONCLUSION & FUTURE WORK

6.1 Findings

- **Ensemble Model Superiority:** Ensemble models—specifically Random Forest and XGBoost—outperformed traditional single algorithms such as Logistic Regression and Decision Tree, achieving over 99% accuracy and high recall. This demonstrates their robustness in detecting rare fraudulent transactions in highly imbalanced datasets.
- **Class Imbalance Addressed Effectively:** The introduction of SMOTE for class balancing had a marked improvement on fraud detection metrics. Models trained without balancing showed low recall, but post-SMOTE, both recall and F1-score increased substantially, proving that balancing minority classes is essential.
- **Critical Role of Data Preprocessing:** Steps such as normalization, encoding of categorical features, and careful feature selection collectively enhanced the predictive power and stability of the models. Effective data preprocessing was found to be just as important as the choice of algorithm.
- **User-Centric Interface and Real-Time Results:** The Streamlit-based web application enabled seamless, real-time user interaction, supporting both single and batch predictions. Users benefited from instant feedback, clear result visualization (including confusion matrices and ROC curves), and accessible system controls.
- **Performance, Scalability, and Robustness:** Rigorous testing—covering both system performance under load and robustness to varying data qualities—showed that the final system was reliable, scalable to large transaction volumes, and resilient against overfitting or model drift.
- **Real-World Applicability:** Overall, the approach demonstrated that a well-engineered machine learning system can provide highly accurate, adaptable, and user-friendly solutions for real-world financial fraud detection, providing significant value to banking and payment industries.

6.2 Conclusion

The **FraudLens** project successfully achieved its objective of developing a practical, scalable, and robust solution for identifying fraudulent transactions using machine learning. By carefully preprocessing the highly imbalanced dataset—including the application of SMOTE for class balancing—and selecting relevant features, the models were effectively prepared to distinguish between legitimate and fraudulent activities.

Comparisons among several machine learning algorithms established that ensemble methods, particularly Random Forest and XGBoost, consistently offered superior detection accuracy and recall. These models proved especially competent at uncovering the minority class (fraudulent transactions), which is critically important in the financial sector where missing a fraud can have serious consequences.

The project also demonstrated the value of integrating the trained models into a user friendly Streamlit web application. This interface allowed users to perform both single and batch transaction analyses with ease, receive immediate classification results, and view informative metrics such as accuracy, precision, recall, and confusion matrices. It ensured accessibility for both technical and non-technical users, making complex machine learning accessible to end-users without coding expertise.

Additionally, the modular design of both the data pipeline and application ensures future extensibility. New algorithms, data sources, or features can be smoothly added as fraud tactics evolve. This adaptability positions the system for long-term relevance and improvement.

In summary, the project illustrates that with proper data preparation, model selection, and thoughtful application development, machine learning can play a pivotal role in modern financial fraud detection. The deployed solution meets key requirements for accuracy, speed, usability, and security, and can serve as a foundation for future enhancements as transaction patterns and fraud behaviors continue to change.

6.3 Future Scope

While the current project provides an effective and user-friendly credit card fraud detection system, there are several opportunities to further enhance its capabilities in future work:

6.3.1 Incorporation of Additional Models

Future development can explore more advanced machine learning and deep learning algorithms, such as neural networks or ensemble stacking methods, which may offer even greater predictive power, especially as more data becomes available.

6.3.2 Real-time Streaming Data Support

The present system is optimized for batch and individual transaction analysis. Integrating real-time transaction monitoring using technologies such as Apache Kafka or Spark Streaming could enable immediate fraud detection and intervention as transactions occur in real-world scenarios.

6.3.4 Enhanced Feature Engineering

Further feature engineering, such as temporal analysis (examining sequences of transactions), geographic behavior analysis, or social network features, could provide richer signals for improving model performance and detecting sophisticated fraud tactics.

6.3.5 Integration with External APIs

Linking the system with external fraud intelligence feeds or banking APIs could provide additional context and up-to-date indicators of new fraud schemes, helping the system to detect previously unseen threats.

6.3.6 User Notifications and Automated Action

Future versions could incorporate user notification modules—such as email or SMS alerts for flagged transactions—or even automated blocking and escalation processes based on model outputs.

References

1. "Fed-RD: Privacy-Preserving Federated Learning for Financial Crime Detection," *arXiv preprint arXiv:2408.01609*, 2024.Arslana, M., Ghanema, H., Munawarb, S., & Cruza, C. (2024). A Survey on RAG with LLMs. *arXiv preprint arXiv:2405.08770*.
2. P. Sundaravadivel *et al.*, "Optimizing Credit Card Fraud Detection with Random Forests and SMOTE," *Sci. Rep.*, 2025.Guettala, M., Bourekkache, S., Kazar, O., & Harous, S. (2024). Building advanced RAG Q&A with Multiple Data Sources Using Langchain: A Multi-Search Agent RAG Application in Ubiquitous Learning. *arXiv preprint arXiv:2402.13286*.
3. "What is Incremental Learning? Applications in Real-Time Fraud Detection," DataCamp Blog, 2023.Ho, J., Chan, W., Saharia, C., Whang, J., Gao, R., Gritsenko, A., Kingma, D. P., Poole, B., Norouzi, M., Fleet, D. J., & Salimans, T. (2022). IMAGEN VIDEO: High Definition Video Generation with Diffusion Models. *arXiv preprint arXiv:2210.02303*.
4. A. Dal Pozzolo, G. Boracchi, O. Caelen, C. Alippi, and G. Bontempi, "Credit card fraud detection: a realistic modeling and a novel learning strategy," *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 29, no. 8, pp. 3784–3797, 2015.Karpukhin, V., Oguz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen, D., & Yih, W.-T. (2020). Dense Passage Retrieval for Open-Domain Question Answering. *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing, EMNLP 2020*.
5. J. Jurgovsky, M. Granitzer, K. Ziegler, and S. Calabretto, "Sequence classification for credit-card fraud detection," *IEEE Access*, vol. 6, pp. 28441–28455, 2018.P, S., Jain, V., Golugula, S., Motamarri, S., & Sathvik, S. (2025). Manimator: Transforming Research Papers and Mathematical Concepts into Visual Explanations. *arXiv preprint arXiv:2502.19321*.
6. F. Carcillo, A. Dal Pozzolo, Y.-A. Le Borgne, O. Caelen, Y. Mazzer, and G. Bontempi, "SCARFF: a scalable framework for streaming credit card fraud detection using Random Forests," *arXiv preprint arXiv:1905.05002*, 2019.Swacha, J., & Gracel, M. (2025).

Retrieval-Augmented Generation (RAG) Chatbots for Education: A Survey of Applications. *arXiv preprint arXiv:2504.05318*.

7. A. C. Bahnsen, D. Aouada, and B. Ottersten, "Feature engineering strategies for credit card fraud detection," *IEEE Access*, vol. 4, pp. 5788–5797, 2016. Workshop at the 2nd ACM Workshop in AI-powered Question and Answering Systems (AIQAM '25). (2025). A Survey of Reasoning and Agentic Systems in Time Series with Large Language Models. *Proceedings of the 2nd ACM Workshop in AI-powered Question and Answering Systems*.
8. C. Whitrow, D. J. Hand, P. Juszczak, D. Weston, and N. M. Adams, "Transaction aggregation as a strategy for credit card fraud detection," *J. Oper. Res. Soc.*, vol. 60, no. 12, pp. 1969–1975, 2009.
9. T. Fawcett, "An introduction to ROC analysis," *Pattern Recognit. Lett.*, vol. 27, no. 8, pp. 861–874, 2006.
10. Y. Sahin and E. Duman, "Detecting Credit Card Fraud by Decision Trees and Support Vector Machines," in *Proc. Int. Conf. Computational Intelligence and Security*, 2011.
11. P. Hurlen *et al.*, "Ethical Challenges in Machine Learning for Fraud Detection," *AI Ethics*, 2021.
12. "Credit Card Fraud Detection Using Machine Learning: Addressing Class Imbalance with Resampling Techniques," SSRN Working Paper, 2025.